

Scala

Lightweight Modular Staging

Steven Both, Toby Rufinus, Jaspreet Singh, Daniël
Stekelenburg



Now that we are familiar with Scala...



Now that we are familiar with Scala...

lets look at an awesome library implemented in Scala.



Now that we are familiar with Scala...
lets look at an awesome library implemented in Scala.
Lightweight Modular Staging (LMS)



LMS is..

'A library-based multi-stage programming approach that uses types to distinguish between binding time.'



LMS is..

'A library-based multi-stage programming approach that uses types to distinguish between binding time.'

Lightweight Modular Staging (LMS) is a runtime code generation approach. The framework provides a library of core components for building high performance code generators and embedded compilers in Scala.



Outline

- ▶ A gentle introduction to LMS
- ▶ Generative Programming
- ▶ Intermediate representation
- ▶ How do we program in LMS?



A gentle introduction to LMS

Power function in Scala

```
def power(b: Double, x: Int): Double =  
  if (x == 0) 1.0 else b * power(b, x - 1)
```



A gentle introduction to LMS

Power function in Scala

```
def power(b: Double, x: Int): Double =  
  if (x == 0) 1.0 else b * power(b, x - 1)
```

Power function in Scala LMS

```
trait PowerA { this: Arith =>  
  def power(b: Rep[Double], x: Int): Rep[Double] =  
    if (x == 0) 1.0 else b * power(b, x - 1)  
}
```



A gentle introduction to LMS

What is multi-stage programming (MSP) and Generative Programming?

- ▶ MSP is a form of metaprogramming in which compilation is divided into a series of intermediate phases, allowing typesafe run-time code generation.
- ▶ In LMS the stages are defined by the types that are used: T (compile-time) vs $\text{Rep}[T]$ (run-time).



Generative Programming

Commonalities between LMS and generative programming languages

- ▶ Code generators and generated code are expressed in the same program.
- ▶ Objects that are live within the generator's heap can be accessed from generated code if the code is invoked directly (cross-stage persistence).



Generative Programming

Commonalities between LMS and generative programming languages - Cont.

- ▶ Staged expressions inherit the static scope of the generator and if the generator is well-typed so is the generated code.
- ▶ Data types representing staged expressions are inaccessible to the program itself making optimizations safe that preserve only semantic but not structural equality.



Generative Programming

Differences between LMS and generative programming languages

- ▶ Staging is determined entirely by types, no special syntax is required. Other languages use quasi-quotations:

```
def power(b: String, n: Int): String =  
  if (n == 0) s"1.0"  
  else s"( b * { power(b, n - 1) } )"
```

- ▶ Given a sufficiently expressive programming language, the whole framework can be implemented as a library.



Differences between LMS and generative programming languages - Cont.

- ▶ Staged code fragments are composed through explicit operations.
- ▶ The relative evaluation order of expressions is preserved across stage boundaries. There is no danger of accidentally omitting, reordering, or duplicating computations.



Drawbacks of LMS

- ▶ The Scala LMS library only implements staged operations for a subset of Scala.
- ▶ Debugging is painful - LMS can give obscure errors, run into an infinite loop, or generate wrong code.
- ▶ Documentation is lacking.



Productivity vs Performance

Software performance depends more on programmer productivity

- ▶ Processor clock speed doesn't double every 18 months
- ▶ High-level programming is hard to translate to efficient code
- ▶ Shift towards big data workloads

Result: Hand-optimized programs (BAD IDEA)

- ▶ abandoning all best practices and benefits of high-level programming
- ▶ programs become hard to read, maintain, verify...
- ▶ this attracts bugs and security vulnerabilities



Solution: Generative Programming

Write a program generator

- ▶ Produces the code of a program as output
- ▶ Reorganization of a programs' execution into stages, also called multi-stage programming



How do we write MSP-programs?

- ▶ A single-stage program is developed, implemented and tested
- ▶ Ensure the the program can be used in a staged manner. Otherwise “refactor”
- ▶ Introduce staging annotations



Another stage!

From

- ▶ Compilation-based program execution:
- ▶ Compile-time, run-time.

To

- ▶ Generated program execution:
- ▶ Generation-time, Compile-time, run-time.



Intermediate representation (IR)

Trees

- ▶ The goal is to get a representation of the implementation which supports optimizing compilation.
- ▶ A common approach is the use of expression trees which resemble abstract syntax trees (AST) with the following types:

```
type Exp[T]           // atomic:   Sym, Const
type Def[T]           // composite: Exp + Exp, Exp * Exp, ...
type Stm[T]           // statement: val x = Def
type Block[T]         // blocks:   { Stm; ...; Stm; Exp }
```



Intermediate representation

Trees

- ▶ Code is generated from the tree by using Forward Traversal
- ▶ Optimizations are applied by applying transformations on the tree.



Intermediate representation

Trees

- ▶ Code is generated from the tree by using Forward Traversal
- ▶ Optimizations are applied by applying transformations on the tree.
- ▶ PROBLEM: Phase Ordering: 'How to determine the ideal sequence of optimization phases to apply to each function or program so as to maximize the gain in speed, code-size, power, or any combination of these performance constraints.'
- ▶ Different orderings can have major impact on performance, especially in performance-critical applications.



Intermediate representation

Graphs

- ▶ Directed Graphs “Sea of nodes”
- ▶ Two types of nodes:
 1. Expressions: constants and symbols
 2. Definitions: composite operations
- ▶ Each definition has at least one associated symbol (expression) and refers to other definitions by the symbols the other definitions are associated with.



Intermediate representation

Graphs

If we consider a purely functional language subset: Graphs..

- ▶ Allow for possibilities for aggressive optimizations;
- ▶ Make optimizations easier to implement;
- ▶ Mitigate the phase ordering problem.
- ▶ Because we rely on referential transparency: the value of an expression is always the same, no matter when and where it is computed, so optimizations do not need to check availability or lifetimes of expressions.



Graphs - Optimizations

- ▶ Global common subexpression elimination (CSE): elimination of identical expressions (evaluate to the same value);
- ▶ Dead code elimination (DCE): finding all reachable statements and discarding everything else.
- ▶ We can define our own optimizations, both generic optimizations and domain-specific ones.
- ▶ Many other optimizations implemented



Intermediate representation

Graphs - Back to code

- ▶ Code motion algorithm
- ▶ The algorithm will try to push statements inside conditional branches and lift statements out of loops.
- ▶ Code motion depends on dependency and frequency information but not directly on data-flow information.



How to LMS:

- ▶ Staging
- ▶ Generating code
- ▶ Data types



How to LMS: Simple staging

Power function in Scala:

```
def power(b: Double, p: Int): Double = {  
  if (p == 0)  
    1.0  
  else  
    b * power(b, p - 1)  
}
```



How to LMS: Simple staging

Staged power function in LMS:

```
def power(b: Rep[Double], p: Int): Rep[Double] = {  
  if (p == 0)  
    1.0  
  else  
    b * power(b, p - 1)  
}
```



How to LMS: Simple staging

Staged power function in LMS:

```
def power(b: Rep[Double], p: Int): Rep[Double] = {  
  if (p == 0)  
    1.0  
  else  
    b * power(b, p - 1)  
}
```

```
power(b, 3)
```



How to LMS: Simple staging

```
power(b, 3)
```

Generated code:

```
def apply(x3: Double): Double = {  
    val x4 = x3 * x3  
    val x5 = x3 * x4  
    x5  
}
```



How to LMS: Staging with recursion

Simple factorial function

```
def fac(n: Rep[Int]): Rep[Int] = {  
  if (n == 0) 1  
  else n * fac(n - 1)  
}
```



How to LMS: Staging with recursion

Simple factorial function

```
def fac(n: Rep[Int]): Rep[Int] = {  
  if (n == 0) 1  
  else n * fac(n - 1)  
}
```

fac(n)



How to LMS: Staging with recursion

Simple factorial function

```
def fac(n: Rep[Int]): Rep[Int] = {  
  if (n == 0) 1  
  else n * fac(n - 1)  
}
```

```
fac(n)
```

```
...
```

```
[error] (run-main) java.lang.StackOverflowError
```

```
[error] (compile:run) Nonzero exit code: 1
```

```
...
```



How to LMS: Staging with recursion

`power(b, 3)` vs. `fac(n)`



How to LMS: Staging with recursion

`power(b, 3)` vs. `fac(n)`

Make use of a lambda function:

```
def fac: Rep[Int => Int] = doLambda { n =>
  if (n == 0) 1
  else n * fac(n-1)
}
```

Now we can try it again:

`fac(n)`



How to LMS: Staging with recursion

Generated code

```
def apply(x12:Int): Int = {  
  var x1 = {x2: (Int) =>  
    val x3 = x2 == 0  
    val x8 = if (x3) { 1 }  
    else {  
      val x4 = x2 - 1  
      val x5 = x1(x4)           // recursion  
      val x6 = x2 * x5  
      x6 }  
    x8: Int }  
  val x13 = x1(x12)           // recursion  
  x13 }
```



How to LMS: Generating code

```
power(b, 3)
```

Generated code:

```
def apply(x3: Double): Double = {  
    val x4 = x3 * x3  
    val x5 = x3 * x4  
    x5  
}
```



How to LMS: Generating code

Optimized power function

```
def powerOpt(b: Rep[Double], p: Int): Rep[Double] = {  
  def loop(x: Rep[Double], ac: Rep[Double],  
           y: Int): Rep[Double] =  
    if (y == 0)  
      ac  
    else if (y % 2 == 0)  
      loop(x * x, ac, y / 2)  
    else  
      loop(x, ac * x, y - 1)  
  
  loop(b, 1.0, p)  
}
```



How to LMS: Generating code

```
powerOpt(b, 3)
```

Generated code:

```
def apply(x3: Double): Double = {  
    val x4 = x3 * x3  
    val x5 = x3 * x4  
    x5  
}
```



How to LMS: Generating code

```
powerOpt(b, 3)
```

Generated code:

```
def apply(x3: Double): Double = {  
    val x4 = x3 * x3  
    val x5 = x3 * x4  
    x5  
}
```

LMS can generate the same code from different staged codes.



How to LMS: Generating code

But not per se.



Universiteit Utrecht

[Faculty of Science
Information and Computing
Sciences]

How to LMS: Generating code

But not per se.

For example:

```
powerOpt(b, 6)
```



How to LMS: Generating code

But not per se.

For example:

```
powerOpt(b, 6)
```

Generated code of optimized version

```
def apply(x4:Double): Double = {  
    val x5 = x4 * x4  
    val x6 = x5 * x5  
    val x7 = x5 * x6  
    x7  
}
```



How to LMS: Generated code

Trivial regular expression

```
checkRegex(".", "Hello world")
```

The char '.' is a wildcard in Scala.



How to LMS: Generated code

Partial regex code

```
def matchStar(...): Rep[Boolean] = { ... }  
def matchBegin(...): Rep[Boolean] = { ... }  
def matchEnd(...): Rep[Boolean] = { ... }  
  
def matchChar(c: Char, t: Rep[Char]): Rep[Boolean] =  
{ c == '.' || c == t }
```



How to LMS: Generated code

Partial regex code

```
def matchStar(...): Rep[Boolean] = { ... }  
def matchBegin(...): Rep[Boolean] = { ... }  
def matchEnd(...): Rep[Boolean] = { ... }  
  
def matchChar(c: Char, t: Rep[Char]): Rep[Boolean] =  
{ c == '.' || c == t }
```

Now do:

```
checkRegex(".", s)
```



How to LMS: Generated code

```
...  
val x47 = while ({ val x28 = x27  
  val x34 = if (x28) { false }  
  else { val x30 = x26  
    val x32 = x30 < x31  
    x32 } x34}) {  
  val x36 = x26 += 1  
  val x37 = x26  
  val x38 = x37 < x31  
  val x42 = if (x38) {  
    val x39 = x25(x37)  
    val x40 = '.' == x39 // matchChar(...)  
    val x41 = true || x40 // c == '.' || c == t  
    x41 }  
  else { false }
```



How to LMS: Generated code

Generated code is not meant to be human-readable



How to LMS: Data types

The previous examples only considered:

- ▶ `Rep[Int]`
- ▶ `Rep[Double]`
- ▶ `Rep[Char]`
- ▶ `Rep[Boolean]`



How to LMS: Data types

The previous examples only considered:

- ▶ Rep[Int]
- ▶ Rep[Double]
- ▶ Rep[Char]
- ▶ Rep[Boolean]

But what to do for datatypes that are of your own making?



How to LMS: Data types

- ▶ Scala LMS only has pre-defined operations for standard library types.
- ▶ You add `Rep[T]` types to your own functions to stage them, allowing them to work on staged values.



How to LMS: Data types

Unstaged:

```
case class Vec3(x: Double, y: Double, z: Double) {  
  def +(that: Vec3): Vec3 =  
    Vec3(this.x + that.x,  
          this.y + that.y, this.z + that.z)  
}
```



How to LMS: Data types

Unstaged:

```
case class Vec3(x: Double, y: Double, z: Double) {  
  def +(that: Vec3): Vec3 =  
    Vec3(this.x + that.x,  
          this.y + that.y, this.z + that.z)  
}
```

Staged:

```
case class Vec3(x: Rep[Double], y: Rep[Double],  
                z: Rep[Double]) {  
  def +(that: Vec3): Vec3 =  
    Vec3(this.x + that.x,  
          this.y + that.y, this.z + that.z)  
}
```



How to LMS: Data types

What about this function?

Unstaged:

```
case class Vec3(x: Double, y: Double, z: Double) {  
  def length: Double =  
    sqrt(x * x + y * y + z * z)  
}
```



How to LMS: Data types

What about this function?

Unstaged:

```
case class Vec3(x: Double, y: Double, z: Double) {  
  def length: Double =  
    sqrt(x * x + y * y + z * z)  
}
```

Staged:

```
case class Vec3(x: Rep[Double], y: Rep[Double],  
  z: Rep[Double]) {  
  def length: Double =  
    sqrt(x * x + y * y + z * z)  
}
```



A problem: undefined operations on staged types

```
[error] main.scala:9: not found: value sqrt
[error]       sqrt(x * x + y * y + z * z)
[error]       ^
[error] one error found
[error] (compile:compileIncremental)
Compilation failed
```

- ▶ Scala LMS defines some fundamental operations on staged types, such as integer/floating point arithmetic.
- ▶ If your staged function only uses those operations (like the + operator in the Vec3 example) you're fine.



A problem: undefined operations on staged types

- ▶ But LMS doesn't define everything in the Scala language! And it certainly doesn't define any third-party library functions.
- ▶ Some staged operations can't be defined in terms of other staged functions. We really don't want to implement `sqrt()` ourselves; we want to use the `scala.math.sqrt()` function somehow.
- ▶ That function call needs to end up in the generated code.
- ▶ We need to be able to add new operations on staged types.



How to Stage Your Algorithm

1. Add `Rep [T]` type annotations
2. Define an interface for new operations on staged types
3. Implement the interface in terms of IR nodes
4. (optional) Define optimizations, rewriting certain patterns of IR nodes
5. Extend code generator so new IR nodes can be turned into code



An example: Fast Fourier Transform

You can create a wave-like function by summing up a number of sine waves. The Fourier transform decomposes the waveform back into its sine components. The Fast Fourier Transform (FFT) is a fast numerical algorithm that can do this.



An example: Fast Fourier Transform

```
def fft(xs: Array[Complex]): Array[Complex]
= xs match {
  case (x :: Nil) => xs
  case _ =>
    val N = xs.length // assume it's a power of two
    val (even0, odd0) = splitEvenOdd(xs)
    val (even1, odd1) = (fft(even0), fft(odd0))
    val (even2, odd2) = (even1 zip odd1 zipWithIndex)
      case ((x, y), k) =>
        val z = omega(k, N) * y
        (x + z, x - z)
    }.unzip;
  even2 ::: odd2
}
```



An example: Fast Fourier Transform

```
case class Complex(re: Double, im: Double) {  
  def +(that: Complex) = Complex(this.re + that.re,  
    this.im + that.im)  
  def *(that: Complex) = ...  
}  
  
def omega(k: Int, N: Int): Complex = {  
  val kth = -2.0 * k * Math.Pi / N  
  Complex(cos(kth), sin(kth))  
}
```

Note the operations we perform on doubles: arithmetic (addition, multiplication, ...) and trigonometry (sin, cos)



An example: Fast Fourier Transform

Step 1: Add Rep[T] type annotations

Unstaged:

```
case class Complex(re: Double, im: Double) {  
  def +(that: Complex) = Complex(this.re +  
    that.re, this.im + that.im)  
  def *(that: Complex) = ...  
}  
def omega ...  
def fftt ...
```



An example: Fast Fourier Transform

Step 1: Add Rep[T] type annotations

Staged:

```
trait FFT { this: Arith with Trig =>
  case class Complex(re: Rep[Double], im:
    Rep[Double]) {
    def + Complex(this.re + that.re,
      this.im + that.im)
    def * ...
  }
  def omega ...
  def fft ...
}
```



An example: Fast Fourier Transform

Step 1: Add Rep[T] type annotations

```
trait FFT { this: Arith with Trig =>
  case class Complex(re: Rep[Double], im:
    Rep[Double]) {
    def + Complex(this.re + that.re,
      this.im + that.im)
    def * ...
  }
  def omega ...
  def fft ...
}
```



An example: Fast Fourier Transform

Step 1: Add `Rep[T]` type annotations

The staged version of complex numbers consists of a pair of `Rep[Double]`. We need to be able to do arithmetic and trigonometric operations on staged doubles. Fortunately, the Scala LMS library happens to define operations for `Rep[Double]`.



An example: Fast Fourier Transform

Step 1: Add Rep[T] type annotations

```
trait FFT { this: Arith with Trig =>
  case class Complex(re: Rep[Double], im:
    Rep[Double]) {
    def + Complex(this.re + that.re,
      this.im + that.im)
    def * ...
  }
  def omega ...
  def fft ...
}
```



An example: Fast Fourier Transform

Step 1: Add `Rep[T]` type annotations

As an exercise, let's pretend that LMS didn't have definitions for `Rep[Double]`. What if we need to define these staged operations ourselves? (In other words: how is `Rep[Double]` implemented in the LMS library?) If we know how to do this, we can also define staged operations on our own data types.



An example: Fast Fourier Transform

Step 1: Add Rep[T] type annotations

```
trait FFT { this: Arith with Trig =>
  case class Complex(re: Rep[Double],
    im: Rep[Double]) {
    def + Complex(this.re + that.re,
      this.im + that.im)
    def * ...
  }
  def omega ...
  def fft ...
}
```



- ▶ How to define our own staged operations that work on `Rep[Double]`?
- ▶ We define these operations in traits and mix them in.
- ▶ The `this: Arith with Trig` part means: whenever the trait `FFT` is instantiated, we need to mix in traits that provide arithmetic and trigonometric operations too.



An example: Fast Fourier Transform

Step 2: Define an interface for new operations on staged types

Scala LMS provides a Base trait; the Rep[T] type is defined there.

```
trait Arith extends Base {  
  def infix_+(x: Rep[Double], y:  
    Rep[Double]): Rep[Double]  
  def infix_*(x: Rep[Double], y:  
    Rep[Double]): Rep[Double]  
  ...  
}
```

```
trait Trig extends Base {  
  def cos(x: Rep[Double]): Rep[Double]  
  def sin(x: Rep[Double]): Rep[Double]
```

[Faculty of Science
Information and Computing
Sciences]



An example: Fast Fourier Transform

Step 2: Define an interface for new operations on staged types

These traits contain only abstract members; they are interfaces. We need to create subclasses with concrete implementations.



An example: Fast Fourier Transform

Step 3: Implement the interface in terms of IR nodes

- ▶ Scala LMS uses a node-based intermediate representation for staged expressions.
- ▶ The `BaseExp` class from the LMS framework defines some related types.
- ▶ `Exp[T]` is a simple IR expression (a constant or symbol).
- ▶ `Def[T]` is a composite operation; these operations will be converted to simple expressions.
- ▶ `BaseExp` also defines `Rep[T] = Exp[T]` so staged expressions will be converted to IR expressions.



An example: Fast Fourier Transform

Step 3: Implement the interface in terms of IR nodes

```
trait ArithExp extends Arith with BaseExp {  
  // These case classes are IR nodes  
  case class Plus(x: Exp[Double], y: Exp[Double])  
    extends Def[Double]  
  case class Times(x: Exp[Double], y: Exp[Double])  
    extends Def[Double]  
  
  // The abstract functions defined in trait Arith are  
  def infix_+(x: Exp[Double], y: Exp[Double]) =  
    Plus(x, y)  
  def infix_*(x: Exp[Double], y: Exp[Double]) =  
    Times(x, y)  
}
```



An example: Fast Fourier Transform

Step 3: Implement the interface in terms of IR nodes

```
trait TrigExp extends Trig with BaseExp {  
  case class Sin(x: Exp[Double]) extends Def[Double]  
  case class Cos(x: Exp[Double]) extends Def[Double]  
  
  def sin(x: Exp[Double]) = Sin(x)  
  def cos(x: Exp[Double]) = Cos(x)  
}
```



An example: Fast Fourier Transform

Step 3: Implement the interface in terms of IR nodes

`sin(x + 2 * y) + sin(0)`

`Plus(Sin(Plus(Sym(x),
Times(Const(2),
Sym(y)))),
Sin(Const(0)))`



An example: Fast Fourier Transform

Step 4: Define optimizations, rewriting certain patterns of IR nodes

- ▶ Scala LMS already contains a number of generic optimizations, such as dead code elimination and reusing identical expressions.
- ▶ We can define our own optimizations, both generic optimizations and domain-specific ones.
- ▶ These are again defined in traits, so you can combine them in a modular way.



```

trait ArithExpOpt extends ArithExp {
  override def infix_*(x: Exp[Double],
    y: Exp[Double]) =
    (x, y) match {
      // Multiplying two constants? We can calculate it
      case (Const(x), Const(y)) => Const(x * y)
      // 1 * x = x, and vice versa
      case (x, Const(1)) => x
      case (Const(1), x) => x
      // Base case: apply the regular base function
      case _ => super.infix_*(x, y)
    }
}

```



An example: Fast Fourier Transform

Step 4: Define optimizations, rewriting certain patterns of IR nodes

```
trait TrigExpOptFFT extends TrigExpOpt {  
  override def cos(x: Exp[Double]) = x match {  
    case Const(x)  
      if { val z = x / math.Pi / 0.5;  
          z != 0 && z == z.toInt } =>  
        Const(0.0)  
    case _ => super.cos(x)  
  }  
}
```



An example: Fast Fourier Transform

Step 5: Extend code generator so new IR nodes can be turned into code

Finally, IR nodes have to be converted to actual code. The LMS framework provides a `ScalaGenBase` class that we can use. We only have to define what to do when the generator encounters one of the new nodes we added.



An example: Fast Fourier Transform

Step 5: Extend code generator so new IR nodes can be turned into code

```
trait ScalaGenArith extends ScalaGenBase
  with ArithExp {
    override def emitNode(sym: Sym[T],
      node: Def[T]) =
      node match {
        // val z = x + y
        case Plus(x, y) =>
          println("val %s = %x + %y".format(sym, x, y))
          // val z = x * y
        case Times(x, y) =>
          println("val %s = %x * %y".format(sym, x, y))
        case _ => super.emitNode(sym, node)
      }
  }
```



An example: Fast Fourier Transform

Step 5: Extend code generator so new IR nodes can be turned into code

- ▶ It's also possible to generate code for other languages if you define your own generator.
- ▶ Scala LMS even comes with a CGenBase trait that can generate C code from your staged Scala functions.



An example: Fast Fourier Transform

Step 5: Extend code generator so new IR nodes can be turned into code

The `CompileScala` trait defines a `compile` function that lets you load the generated code immediately into the running program. Essentially, `compile` “unstages” your staged function (`Rep[A] => Rep[B]`) into a regular function (`A => B`). This function can then be called:

```
val fftCompiled = compile(fft)
// Now we can call fftCompiled with regular values
// Just like any other function in the program
fftCompiled(Array(1.0,0.0, 1.0,0.0, 2.0,0.0, 2.0,0.0))

> Array(6.0,0.0, -1.0,1.0, 0.0,0.0, -1.0,-1.0)
```



Summary

- ▶ LMS is a powerful tool in creating high performance DSL's.
- ▶ Staging based on types makes it different from current approaches.
- ▶ Programs are hard to debug
- ▶ The intermediate representation allows for aggressive optimizations.
- ▶ We can define our own optimizations, both generic optimizations and domain-specific ones.

